

Jetson Model Zoo — Baseline Results

27 models × 12 prompts × 2048 max tokens | temp=0.3, top_k=40, top_p=0.9, repeat=1.1 | GPU (-ngl 999, -fa on)

Rank	Model	Size	HTML	Python	Poetry	Math	Creative	Func Call	AVG	t/s
1	hermes3-3b-q5	3-4B	10.0	10.0	8.0	9.0	10.0	6.2	7.8	19.4
2	smallthinker-3b	3-4B	8.5	8.0	3.0	9.0	8.0	7.7	7.6	20.1
3	qwen2.5-3b	3-4B	9.5	9.0	8.0	8.0	9.0	6.2	7.5	20.1
4	granite4.1-3b	3-4B	9.0	10.0	8.0	8.0	8.0	5.8	7.2	18.9
5	ministral-3b-reasoning	3-4B	10.0	10.0	10.0	8.0	9.0	5.0	7.2	18.7
6	hermes3-3b-q4	3-4B	9.5	10.0	5.0	8.0	10.0	5.5	7.1	19.7
7	granite3.2-2b	2-2.6B	9.5	10.0	7.0	9.0	7.0	5.3	7.0	26.6
8	qwen3.5-2b	2-2.6B	9.5	10.0	6.0	8.0	8.0	5.0	6.8	24.9
9	stablelm-zephyr	?	10.0	10.0	8.0	9.0	8.0	4.3	6.8	26.4
10	smollm3	?	8.0	10.0	8.0	8.0	8.0	4.8	6.6	20.7
11	granite4.2-3b	3-4B	10.0	10.0	6.0	8.0	8.0	4.3	6.5	18.8
12	qwen3-1.7b	1.7B	9.5	10.0	2.0	9.0	7.0	5.0	6.4	30.8
13	granite4-3b	3-4B	9.5	10.0	5.0	8.0	7.0	4.3	6.2	18.9
14	qwen2.5-coder-3b	3-4B	10.0	10.0	5.0	8.0	10.0	3.3	6.1	20.0
15	lfn2.5-2.6b	?	10.0	10.0	3.0	8.0	8.0	3.3	5.8	23.8
16	ministral-3b	3-4B	6.5	7.0	2.0	4.0	7.0	3.3	4.4	24.8
17	deepseek-r1-1.5b	1.5B	1.0	1.0	5.0	3.0	6.0	2.3	2.6	32.9
18	gemma2-2b	2-2.6B	1.0	2.0	2.0	3.0	2.0	3.2	2.5	22.3
19	llama3.2-3b	3-4B	1.5	3.0	4.0	1.0	3.0	2.7	2.5	19.7
20	llama3.2-1b	1B	1.5	1.0	4.0	1.0	5.0	2.5	2.4	42.2
21	llama3.2-3b-new	3-4B	1.0	2.0	4.0	2.0	4.0	1.8	2.1	19.7

22	phi4-mini	?	1.0	1.0	1.0	4.0	3.0	2.3	2.1	17.0
23	gemma3n-e2b	2-2.6B	0.5	1.0	3.0	3.0	3.0	2.0	1.9	21.2
24	gemma3-1b	1B	1.5	2.0	0.0	1.0	3.0	2.2	1.8	31.5
25	phi3-3.8b	?	1.0	0.0	2.0	2.0	2.0	2.0	1.7	17.9
26	deepseek-r1-7b	7B	0.0	0.0	2.0	0.0	1.0	1.3	0.9	9.9
27	gemma4-e2b	2-2.6B	0.0	1.0	1.0	0.0	2.0	1.0	0.8	19.8

Key Findings

- **Hermes 3 3B Q5** is the overall champion (7.8/10) — perfect 10s on HTML, Python, and Creative Writing.
- **SmallThinker 3B** (#2, 7.6) is the best thinking model — concise reasoning doesn't eat the output budget.
- **Ministral 3B Reasoning** scored a perfect 10 on iambic pentameter — the only model to do so.
- **Qwen 2.5 3B** (#3, 7.5) is the most well-rounded — no category below 6.0.
- **Function calls** are the weakest category across all models (avg 4.3/10) — small models can't do tool use.
- **HTML and Python** are the strongest categories — most 3B models score 8-10.
- **Thinking models** that ramble (DeepSeek R1 7B: 0.9, Phi-4-mini: 2.1) score far worse than concise thinkers.
- **Gemma E2B variants** are broken on llama.cpp — chat template issues cause garbage output (0.8, 1.9).
- **Llama 3.2** models underperform their size class — likely template/chat format issues.
- **Phi-3 3.8B** is the worst non-broken model (1.7) — outdated architecture, poor instruction following.

Category Winners

Category	Winner	Score	Runner-up	Score
HTML	hermes3-3b-q5	10.0	ministral-3b-reasoning	10.0
Python	hermes3-3b-q5	10.0	granite4.1-3b	10.0
Poetry	ministral-3b-reasoning	10.0	hermes3-3b-q5	8.0
Math	hermes3-3b-q5	9.0	smallthinker-3b	9.0

Creative	hermes3-3b-q5	10.0	hermes3-3b-q4	10.0
Func Call	smallthinker-3b	7.7	hermes3-3b-q5	6.2